

RENDER YOUR FIRST FILM

Your computer draws a short film, frame by frame – moonlit water, a drifting paper boat – and narrates YOUR story with a free offline voice.

no editor • no camera • no AI video tools • \$0

BEFORE YOU START – the one honest caveat

This kit is driven by Claude Code, which needs a PAID Claude account (claude.com). That is the only non-free thing here – every tool the kit installs is free and open-source.

No Claude account? You can still run everything by hand – page 3.

A STARTER SLICE – not the full studio

For experiencing the workflow, learning step by step. The channel's films use the full system (page 4), so your renders will look simpler – that's by design.

THREE STEPS

1 Install Claude Code Desktop

Download it at claude.com/claude-code and sign in with your Claude account. Windows or macOS.

2 Paste this one command



```
Clone or download the GitHub repository  
framesfromcode/videos to my home directory, then  
read 02-first-video/SKILL.md inside it and follow  
it exactly.
```

Claude Code checks your machine, installs the free tools (Python, FFmpeg, a ~340 MB offline voice model), and renders the sample film: "The Paper Boat", 55 seconds, drawn and narrated entirely by your computer.

3 Then tell it YOUR story

A memory. A bedtime story for your kid. Something you wish existed. Claude rewrites the screenplay file and your machine renders it. If your story needs a night train or a city skyline, Claude draws the new element too. Or ask smaller:

- "Draw a night train for my father's story."
- "Let me hear the other narrator voices."
- "End at dawn instead."

WHAT'S ACTUALLY HAPPENING

No editor ever opens. Your whole movie lives in one screenplay file (story.py): scenes, narration, designed pauses. Each line is spoken by Kokoro-82M, a free offline text-to-speech model, measured, and every scene lasts exactly as long as its words – the same audio-first idea the whole channel runs on. Then a Python script computes every pixel with numpy + Pillow and pipes the frames into FFmpeg. Deterministic: the same story and voice always render the same film, frame for frame.

Prefer to do it by hand? No Claude needed:



```
pip install numpy pillow kokoro-onnx
python get_tts.py # one-time voice download
python first_film.py --proxy # fast draft
python first_film.py # full 1080p
```

If something goes wrong

x python is not recognized

reinstall Python with 'Add to PATH' ticked, reopen the terminal

x No module named kokoro_onnx

pip install kokoro-onnx (same Python that runs the script)

x ffmpeg is not recognized

restart the terminal after installing FFmpeg

THIS KIT IS A SLICE OF SOMETHING BIGGER

With this kit you can make as many films as you want – and grow its night world with anything your story needs. The pipeline behind the channel's episodes adds the craft that makes a film feel directed:

- > **the directing hand**

cameras that breathe and tremble, staging, micro-events – the film holds your viewer

- > **a style library**

proven visual languages – no two films look alike

- > **music born from code**

a score that knows your story and lands on your biggest line

- > **an audience compass**

topics from measured search demand, titles tested on live autocomplete – films people were already looking for

- > **mastering & QC**

broadcast loudness, frame checks – sounds like TV

Film is only the first lane – the same pipeline renders explainers, data stories, science visuals: quieter niches, lower barriers. More starter kits are on the way.

All of it is taught free in the episodes. For everyone else: a packaged, ready-to-run version is in the works – everything above, working for you one pasted sentence at a time. It saves you the research hours, nothing more: every lesson was paid for in failed renders, and its proof is the film that brought you here. Grabbed this kit on Payhip? You're already on the list. No income promises – tools, not dreams.

framesfromcode youtube.com/@framesfromcode